GURUDATHA KANDANATHI CHITTEM

864-765-7834 | gurudathakandanathichittem@gmail.com | LinkedIn | GitHub | Portfolio | United States

Professional Summary

AI Engineer specializing in forward-deployed production AI — designing, deploying, and optimizing multi-agent systems and retrieval-augmented generation (RAG) in enterprise environments. Expert in LangGraph, LangChain, MCP, LLM evaluation, and FAISS/Pinecone vector search; full-stack depth in Python, TypeScript/React, FastAPI, and PostgreSQL. Integrates AI with existing products and APIs, working with stakeholders and enterprise clients from requirements to production.

Technical Skills

AI & Generative AI: Agentic AI, Multi-Agent Systems, RAG, LangChain, LangGraph, LLM Evaluation, MCP, Embeddings, Vector Search, FAISS, Pinecone, Prompt Engineering, AWS Bedrock, Claude Code
Programming & Frontend: Python, TypeScript, JavaScript, Java, SQL, C#, React, Tailwind CSS, .NET MAUI, XAML
Backend & Security: Node.js, Express, FastAPI, REST APIs, WebSockets, JWT, RBAC, API Integration
Databases & Cloud: PostgreSQL, MySQL, Cosmos DB, BigQuery, AWS, Azure, GCP
DevOps & Testing: Docker, Kubernetes, GitHub Actions, Git, API Testing, Debugging, Functional Testing

Professional Experience

Enteraxion | AI Engineer | Remote Dec 2025 – Present

- Designed, built, and deployed a multi-agent KYC automation system for enterprise client onboarding (LangGraph, LangChain) with task-specific agents, conditional routing, and traceable execution paths — automating 80% of KYC workflows across 1,000+ monthly applications.
- Cut KYC processing error rates from 15% to 2% through agent validation checks, structured outputs, and human-in-the-loop review.
- Implemented multi-model routing across three LLM providers by task complexity, reducing model inference costs by 30%.
- Integrated 10+ verification services through MCP-based tool calling, with PostgreSQL-backed workflow state and resumable execution for multi-step AI operations.
- Built evaluation pipelines with Braintrust, running regression evals on every release to monitor output quality, hallucination risk, and latency.
- Maintain full audit trails and human-in-the-loop review paths for low-confidence cases, so every automated decision stays explainable and reversible.

MyUI.ai | AI Engineer | Remote Aug 2024 – Dec 2025

- Built Python-based RAG workflows for document ingestion, text chunking, embedding generation, and Pinecone/LlamaIndex vector retrieval, indexing 5,000+ documents for source-grounded LLM responses.
- Reduced hallucinated responses by 45% through retrieval tuning, prompt construction, and source-grounding workflows that improve answer traceability.
- Optimized chunking, embedding, and retrieval configuration to hold p95 end-to-end RAG latency under 2 seconds.
- Created a 500+ question-and-answer evaluation set, run before every release, to benchmark answer quality, grounding, and hallucination rates.
- Developed C#, XAML, and .NET MAUI application interfaces integrating REST APIs, reusable components, authenticated navigation, and role-based access.
- Used Claude Code for codebase exploration, debugging, refactoring, and test generation, reviewing suggested changes before integration.

Adroit Global Solutions | Software Engineer | Bengaluru, India Jun 2021 – Jun 2023

- Developed and maintained high-throughput payment processing microservices using Java 11, Spring Boot, and Spring Cloud handling millions of daily transactions with sub-second response times.
- Designed event-driven architectures using Apache Kafka for asynchronous payment event processing, ensuring exactly-once delivery and fault tolerance across distributed systems.

Selected Technical Projects

SENTINEL AI – Autonomous Software Reliability Platform (Architecture Design) | LangGraph, Docker, Git, Static Analysis

- Designed a proposed six-agent architecture comprising Mission Controller, Codebase Analyst, Failure Hypothesis, Root-Cause, Repair, and Adversarial Verification agents for automated software reliability analysis.
- Specified tool integrations for Git, test runners, static analysis, and Docker sandboxing, with independent verification gates for regression and edge-case testing.
- Defined agent communication and orchestration strategies using LangGraph state transitions, shared context, and conditional task routing to coordinate root-cause analysis and repair planning.

GCP Data Analysis | BigQuery, SQL, Python, JSON

- Queried and transformed nested JSON datasets using BigQuery SQL and Python, handling inconsistent fields and preparing structured data for downstream analysis.
- Validated extracted records, checked data consistency, and organized analysis-ready outputs to support reliable reporting.

Education

Clemson University | MS, Computer Science Aug 2023 – May 2025

Certification

AWS Certified Generative AI Developer – Professional